Pulse Adoption — Discovery Report

Synthesized by Meridian from 10 internal interviews

Research Goal

Understand why internal teams aren't adopting Pulse, our internal analytics platform.

Executive Summary

• Pulse has a confirmed trust crisis driven by three compounding bugs/gaps: a known join bug causing double-counting, a 6-hour (sometimes multi-day) sync lag, and zero data lineage or freshness transparency — together these cause users across every team to reject Pulse numbers before stakeholder meetings and revert to source systems or spreadsheets as their real source of truth.

• Adoption has effectively collapsed into workarounds: eight of ten interviewees have built or rely on a substitute (shared spreadsheets, Zendesk dashboards, Grafana, pandas notebooks, or manual analyst requests) — meaning Pulse is generating shadow infrastructure costs and analyst-queue backlog rather than self-serve value.

• There is a fundamental interface split: non-technical users (Sales, Support, CS, Finance, Ops) need pre-built, role-specific dashboards with plain-language metric names and zero query building, while technical users (Engineering, Marketing Analytics, Product Analysis) need raw SQL access, higher query/timeout limits, and an API — the current middle-ground UI satisfies neither group.

• Despite these barriers, latent demand is strong and conditional: all eight interviewees who described workarounds explicitly stated they would return to Pulse if trust, freshness, and access issues were resolved — this is a fixable adoption problem, not a product-concept rejection.

• Access permission friction compounds everything: multi-day ticket queues block new analysts at onboarding and lock power users out of cross-team data, routing work back to the data team and flooding Priya's queue — the permission model may be a legacy artifact rather than an intentional governance design.

Key Findings

• Query builder too complex for power users and non-technical users alike — but for opposite reasons (confirmed, high confidence): Non-technical users find the query builder overwhelming for simple tasks, while technical users hit a ceiling when they need advanced SQL constructs, higher query limits, or API access — the tool satisfies neither end of the spectrum.

• Spreadsheets, source tools, and custom-built solutions become the real source of truth (confirmed, high confidence): Users across Sales Ops, Support, Marketing, Finance, Engineering, Customer Success, Product Analysis, and Operations bypass Pulse entirely — exporting to spreadsheets, staying in source systems, building their own dashboards, pulling data manually, waiting on the data team, or escaping to notebooks — which become the team's de facto source of truth.

• Slow query performance and timeouts undermine trust and usability (confirmed, high confidence): Full-scale queries take so long to load — or time out entirely — that users abandon Pulse mid-workflow; the lag and timeout ceiling scale with data volume and query complexity, meaning the heaviest users suffer most.

• Slow query performance and timeouts undermine trust and usability (confirmed, high confidence): The manual export-plus-pivot or notebook workflow is faster end-to-end than waiting for a single Pulse query to complete, making workarounds the rational choice.

• Data accuracy mismatches and known bugs destroy trust in Pulse numbers (confirmed, high confidence): When Pulse figures disagree with source-system numbers, users lose confidence in all Pulse data, stop presenting it to stakeholders, and rebuild figures manually to defend in meetings.

• Data accuracy mismatches and known bugs destroy trust in Pulse numbers (confirmed, high confidence): A confirmed join bug causes double-counting in at least one data model; this is a root cause of discrepancies users are noticing and the highest-priority fix according to the data engineer.

• Data accuracy mismatches and known bugs destroy trust in Pulse numbers (confirmed, high confidence): Users suspect discrepancies stem from definitional differences but feel they shouldn't have to investigate — Pulse offers no metric-definition transparency or data lineage, and the burden of proof falls on the tool. Non-technical users additionally can't identify which metric to pick in the first place because terminology is opaque.

• Stale data makes Pulse unusable for operational and time-critical decisions (confirmed, high confidence): Data is stale by design (6-hour sync cycle, multi-day lags during close, or days-old account health signals), making Pulse structurally unfit for intraday, event-driven, customer-facing, or month-end decisions across multiple teams.

Evidence Quotes

• "it's too complex to get a simple pipeline number out of it. I couldn't figure out the query builder at all, all these joins and filters and dimension dropdowns. I just need one number." — Dana Whitfield

• "My whole team basically lives in that shared sheet now, it's our real source of truth, not Pulse." — Dana Whitfield

• "I've literally walked to get coffee and come back and it's still loading." — Dana Whitfield

• "The export plus pivot is faster than waiting on one Pulse query to load." — Dana Whitfield

• "Pulse will say we closed 340 tickets yesterday, Zendesk says 312. I have no idea which is right, and if I can't trust the number I'm not going to put it in front of my VP." — Marcus Lee

Contradictions & Nuance

• There is a genuine tension between what different user types need from Pulse's query interface: non-technical users (Dana Whitfield, Marcus Lee, Grace Kim, Owen Bradley, Lucia Romano, Derek Olsen) want simpler pre-built dashboards, plain-language metric names, clear reconciled outputs, and less query-building complexity, while technical users (Tomás Herrera, Hannah Ortiz, Raj Patel, and implicitly Priya Raman) want more power — raw SQL access, higher query timeout limits, an API endpoint, window functions, and fewer guardrails. These are mutually exclusive interface philosophies; optimizing for one risks alienating the other. Hannah's team abandoned Pulse entirely for a custom Grafana dashboard, and Raj escapes to pandas notebooks, rather than work within the UI constraints. Derek explicitly said the new report builder is 'the opposite of what I need.'

Recommended Next Steps

• Fix the confirmed join bug immediately and audit all affected data models (including the finance revenue model behind Grace's $40K GL mismatch and the marketing campaign table behind Tomás's 18% GA discrepancy) — this is the single highest-leverage trust restore action and should be treated as a P0 bug, not a roadmap item.

• Ship a data freshness timestamp and inline metric definitions (hover-over plain-language descriptions and last-refreshed time) as a fast-follow UI change — this is a low-engineering-cost intervention that directly addresses the 'black box' and 'I can't decode this metric name' barriers cited by six interviewees and can precede any sync infrastructure work.

• Accelerate the sync cadence toward near-real-time (or at minimum reduce the 6-hour cycle) and investigate whether Grace's month-end 2–3 day lag and Owen's account health delay share the same pipeline root cause as the 6-hour sync — resolving staleness unblocks the Support, CS, Finance, and Operations personas who need intraday or same-day data.

• Run a two-track UX initiative in parallel: (a) build curated, pre-built role-specific starter dashboards for non-technical personas (Sales Director, CSM, Support Lead, Ops Coordinator, Finance) with plain-language metric names and no query-building required; (b) expose a raw SQL interface, configurable query timeout ceiling, and API endpoint for technical users (Engineering, Marketing Analytics, Product Analysis) — these tracks are not in conflict and can ship independently.

• Redesign the permissions model with the data team: audit whether per-team scoping is an intentional governance decision or a legacy artifact, and evaluate row-level security or data-sharing agreements as alternatives to ticket-based access requests —

the current model is the direct cause of new-analyst drop-off at onboarding and blocks the cross-team analyses that represent Pulse's highest-value use cases.

• Present the full findings — especially the latent demand signal from all eight conditional adopters — to product and engineering leadership with a prioritized fix sequence: (1) join bug, (2) freshness timestamp + metric definitions, (3) sync cadence, (4) permissions redesign, (5) dual UX tracks — and confirm whether each item has an owner and a timeline on the current roadmap.

Next AI Interview Plan

We've heard from Sales, Finance, Engineering, and Support that Pulse isn't fitting into their daily workflows — I'd love to understand specifically how you use data day-to-day and where Pulse does or doesn't show up.

1. Derek Olsen told us he wants exactly three numbers — bookings, pipeline, attainment — visible the moment he logs in, with no filters or building required. Does that description match what you or others on your team need from a daily data touchpoint?

Why this question: Six non-technical users (Dana, Marcus, Grace, Owen, Lucia, Derek) have now independently described wanting pre-built role-specific views. We need to confirm whether this is a universal non-technical user pattern or specific to certain roles, and whether specific metric sets per role are consistent enough to build toward.

Grounded in: Derek Olsen: 'I just want one simple summary screen of my region. Bookings, pipeline, attainment, the three things I actually care about, right there when I log in. No building, no filters. Just show me.'

2. Derek told us his standard fallback is to Slack an analyst when Pulse fails him — and Priya Raman said she's flooded with exactly those requests. How often does your team route data requests through an analyst instead of self-serving in Pulse, and what triggers that decision?

Why this question: Both the demand side (Derek, Lucia, Dana) and the supply side (Priya) confirm analyst-as-middleman is a systemic pattern. We need to quantify how widespread this is and understand whether this is seen as a cost or accepted as normal.

Grounded in: Derek Olsen: 'I went in, picked the wrong dataset apparently, got a number that looked off, second-guessed it, then just Slacked an analyst and asked her to send me the real figure. That's what I do every time, honestly.'

3. Multiple people have told us they picked a dataset in Pulse, got a number that looked wrong, and immediately lost confidence — Grace Kim described a $40K revenue discrepancy, Marcus Lee saw ticket counts that didn't match Zendesk. Has that happened to you, and how did it change how much you rely on Pulse afterward?

Why this question: Data trust is a confirmed high-confidence theme. We want to probe whether the trust collapse is universal once a discrepancy is seen, or whether some users find ways to verify and recover trust. Derek's quote about second-guessing a number fits this pattern for a Sales Director persona we haven't fully explored.

Grounded in: Derek Olsen: 'I went in, picked the wrong dataset apparently, got a number that looked off, second-guessed it.'

4. Derek said he's on the road constantly and is never at his desk when he needs a number. Owen Bradley similarly needs real-time signals from wherever he is. How are field-facing roles like yours currently getting data when you're not at a desktop — and would a lightweight or mobile-accessible Pulse view actually change your behavior?

Why this question: Derek introduced a mobile/field-access dimension not previously raised. We need to test whether this is a widespread constraint for Sales Directors, CSMs, and other field roles — and whether it represents a distinct adoption barrier or is subsumed by the pre-built dashboard need.

Grounded in: Derek Olsen: 'it'd be great as a mobile app, by the way, I'm on the road constantly and I'm never at my desk when I need a number.'

5. Derek described needing his West region bookings for a QBR and not trusting what Pulse returned — ending up using analyst-provided data instead. How do you currently prepare data for leadership-facing meetings like QBRs, and what would it take for Pulse to be the source you'd actually cite in that room?

Why this question: QBR and leadership-presentation scenarios are high-stakes moments where trust failures are most costly. Several interviewees (Tomás, Grace, Marcus) mentioned avoiding Pulse for anything going to leadership. The Sales Director QBR context is a new high-stakes scenario worth probing to understand the 'presentable data' standard.

Grounded in: Derek Olsen: 'Last Monday I needed my West region bookings for the QBR. I went in, picked the wrong dataset apparently, got a number that looked off, second-guessed it, then just Slacked an analyst.'

Open Questions

• The join bug is the confirmed root cause of at least some discrepancies — which specific data models are affected, and does the finance revenue model (source of Grace's $40K GL mismatch) contain the same bug as the marketing campaign table?

• Is there a roadmap or timeline to fix the join bug, and do product/engineering stakeholders understand it as the top trust-blocker across Finance, Marketing, and Support?

• Tomás, Hannah, and Raj all want raw SQL / higher query limits / API access — how many other technical users are hitting the query-builder ceiling or timeout wall, and is a SQL escape-hatch, configurable timeout ceiling, or API endpoint already on the product roadmap?

• Hannah's team built a Grafana dashboard off raw tables and Raj escapes to pandas notebooks — how many other engineering or technical teams have similarly routed around Pulse entirely? Is this shadow-infrastructure pattern tracked anywhere?

• Hannah cited performance scaling with query size and Raj reports outright timeouts on multi-table joins — is the performance bottleneck in the query engine layer, the data warehouse, or the Pulse application layer? Has engineering investigated this?

• Would exposing metric definitions, calculation logic, and data lineage inline (e.g., hover-over tooltips, a plain-language data dictionary, or journal-entry traceability for Finance) reduce the 'I can't defend this number' and 'I can't decode this metric name' barriers enough to keep both technical and non-technical users in Pulse?

Methodology

• Meridian conducted 10 voice interviews with internal Pulse users across teams.

• Each transcript updated a living insight document before the next interview plan was generated.

• Findings confirmed by multiple interviewees were promoted in confidence; one-off gripes were treated as noise.